
1. Spatial Management

Introduction

Spatial management is crucial for making games scale well. Most games take place in spatial environments (2D/3D maps), requiring efficient handling of spatial queries and data structures.

Spatial Management in Games

Key Requirements

- **Transition models, agent control, and transmission protocols** require spatial query processing
- Main spatial queries needed:
 - Which game entities are within interaction range? (Area of Interest - AoI)
 - Support collision detection and area intersections (pre-filtering)
 - Which game entity is closest?
 - Does a player enter the aggro-range of an opponent?

Performance Challenge Example

For small game worlds with limited entities, spatial positions can be organized in a list with sequential scan processing. However, this becomes expensive with high query frequencies and large numbers of moving objects.

Example calculation:

- 1000 game entities in one zone
- 24 ticks/second
- Naive AoI computation requires: $(1000 \times 999)/2 \times 24 = \mathbf{11,988,000 \text{ distance computations per second}}$

Conclusion: The cost for spatial query processing strongly increases with the size of the game state (number of potential interacting entities grows quadratically).

Types of Spatial Queries

1. ϵ -Range Queries

$$RQ(q, \epsilon) = \{v \in GS \mid \sqrt{((q_1 - v_1)^2 + (q_2 - v_2)^2)} \leq \epsilon\}$$

Returns all objects within Euclidean distance ϵ from query point q .

2. Box Query

$$BQ(q, \epsilon) = \{v \in GS \mid x_1 \leq v_1 \leq y_1 \wedge x_2 \leq v_2 \leq y_2\}$$

Returns all objects within a rectangular region.

3. Intersection Query

$$IQ(q, r) = \{(v, s) \in GS \times \mathbb{R} \mid \sqrt{((q_1 - v_1)^2 + (q_2 - v_2)^2)} \leq r + s\}$$

Finds all objects whose regions intersect with the query region.

4. Nearest Neighbor Query

$$NN(q) = \{v \in GS \mid \forall x \in GS: \sqrt{((q_1 - v_1)^2 + (q_2 - v_2)^2)} \leq \sqrt{((q_1 - x_1)^2 + (q_2 - x_2)^2)}\}$$

Finds the closest object to the query point.

Efficiency Tuning Methods

1. Methods to Reduce Number of Considered Objects (Pruning)

- Distribute the game world (zoning, instancing, sharding)
- Index structures (BSP-Tree, KD-Tree, R-Tree, Ball-Tree)

2. Reduce Number of Spatial Queries

- Reduce query ticks
- Spatial publish-subscribe

3. Efficient Query Processing

- Nearest-neighbor queries
- ϵ -range Join (simultaneously compute all Aols)

Spatial Management Techniques

Sharding and Instantiation

- **Concept:** Copying a region for a specific group
- **Benefits:**

- Any number of the same region can exist
- Part of game design (e.g., limiting players for a quest)
- Reduces performance issues in open world
- **Complications:**
 - Doesn't solve underlying problem (no connected game world)
 - Instance management can cause additional expenses
 - Worst case: 1000 parallel game states for 1000 players

Zoning

- **Concept:** Splitting open world into several fixed areas
- **Benefits:**
 - Only objects in current zone need consideration
 - Partitions both space and game state
 - Easier distribution across multiple computers
- **Problems:**
 - Objects of adjacent zones need consideration
 - Heterogeneous distribution of players over zones

Micro-Zoning

- **Concept:** Game world partitioned into small areas (micro zones)
- **Process:**
 - Only entities within actual micro zone are managed
 - Only micro zones intersecting AoI are relevant
 - Sequential search within a region
 - Zones created via grids, Voronoi-cells, etc.

Spatial Publish-Subscribe

- **Concept:** Combination of micro-zoning and subscriber systems
- **Process:**
 1. Game entities register in current micro zone (publish)
 2. Entities subscribe to all micro zones with intersecting Aols
 3. List of entities within AoI created by merging subscribed zones
- **Advantages:**
 - Close-by objects determined efficiently
 - Changes actively sent to subscribers (no regular queries needed)

- **Disadvantages:**
 - Micro zones can become overcrowded
 - Zone changing overhead increases if zones too small
 - Zone border placement may cause extreme fluctuations
 - High change rates increase overhead significantly

Index Structures

Classic Index Structures Overview

- **Space Partitioning:** Partition data space via dimensional splits
 - Examples: Quad-Tree, BSP-Trees
 - May include dead space
 - Potentially bad search performance for spatial queries
- **Data Partitioning:** Page regions defined by minimum bounding regions
 - Example: R-Trees
 - Better pruning performance
 - Page regions may overlap
 - Risk of degeneration

Important Features of Search Trees

- **Page region:** Surrounding approximation of several objects
- **Balancing:** Addresses variance of path lengths or objects per region
- **Page capacity:** Min/max number of objects within page region
- **Overlap:** Intersections between page regions
- **Dead space:** Space without page regions
- **Pruning:** Exclusion of all objects within a page region

Requirements for Games

- Whole tree stored in main memory
- High volatility (position changes every tick)
- Trees might degenerate requiring costly balancing
- Many queries per time unit
- Support for multiple queries during one tick
- Objects have 2-3 dimensions with volume (hitbox)

Key conclusion: Runtime increase for query processing must compensate for time spent on index creation/update.

Binary Space Partitioning (BSP) Trees

kD-tree (most popular BSP variant)

- Root contains whole data space
- Every inner node has two successors
- Data objects stored in leaf nodes
- Properties:
 - Max page capacity: M entries
 - Min page capacity: $M/2$ entries
 - At overflow: splitting w.r.t. an axis
 - Axis changes after every split
 - Data distributed 50%-50%
 - At deletion: merge sibling nodes

Problems with Dynamic Behavior

- No balancing (tree might degenerate)
- Rebalancing possible but expensive
- High update complexity

Bulk-Load for BSP Trees

- Assumption: All data objects known
- Creation: Recursively distribute objects with 50/50 split
- Always creates balanced tree
- Data page of tree height h with n objects contains:
 - At least $\lfloor n/2^h \rfloor$ objects
 - At most $\lfloor n/2^h \rfloor + 1$ objects

Quad-Trees

- Root represents whole data space
- Every inner node has four successors
- Sibling nodes split parent space in four equal parts
- Usually not balanced
- Pages have maximum filling ratio M , no minimum

- Leaves contain data objects

R-Trees

Structure

- Root encompasses complete data space (max M entries)
- Page regions modeled by Minimal Bounding Rectangles (MBR)
- Inner nodes have m to M successors ($m \leq M/2$)
- MBR of successor completely contained within predecessor's MBR
- All leaves at same height
- Leaves contain data objects (points or rectangles)

Insertion Algorithm

Three cases when inserting object x:

1. **x contained in directory rectangle D:** Insert into subtree of D
2. **x contained in several rectangles $D_1 \dots D_n$:** Insert into subtree with smallest area
3. **x not contained in any rectangle:** Insert into subtree with smallest area increase

Split Algorithm (Basic $O(n^2)$)

When node K overflows ($|K| = M+1$):

1. Choose pair (R_1, R_2) with largest "dead space":

$$d(R_1, R_2) = \text{area}(\text{MBR}(R_1 \cup R_2)) - \text{area}(R_1) - \text{area}(R_2)$$

2. Set $K_1 = \{R_1\}$ and $K_2 = \{R_2\}$
3. Repeat until all R_i assigned:
 - If all remaining needed to fill smaller node: assign all and stop
 - Otherwise: allocate next R_i to node with smallest MBR area increase

Linear Split Algorithm $O(n)$

- Same as basic but different initial pair selection
- Choose pair with "greatest distance":
 1. Find rectangles with lowest max and highest min values per dimension
 2. Normalize maximum distance by dividing by total expansion
 3. Choose pair with greatest normalized distance
- Linear complexity: $O(2m+1 \times d)$ where d = dimensions

Bulk-Load and R-Trees

- **Advantages:** Faster creation, better query processing structure
- **Optimization criteria:**
 - Greatest possible filling ratio (low height)
 - Little overlap
 - Small dead space

Sort-Tile-Recursive Algorithm

1. Consider n points/rectangles to store
2. Calculate quantile: $q = \sqrt{(n/M)}$
3. Sort data in dimension 1
4. Generate quantile after $q \cdot M$ objects
5. Sort objects of every quantile in dimension 2
6. Generate quantile after M objects
7. Create MBR around points in each cell
8. Restart with derived MBRs or stop if $q < 2$

- **Time complexity:** $O(n \log n)$
- **Note:** Creates non-overlapping MBRs for points

Deletion in R-Trees

- Test for underflow after deletion: $|S| < m$
- If no underflow: delete and stop
- If underflow:
 1. Delete underflowed page from predecessor
 2. Insert remaining elements into R-Tree
 3. If root has single child, child becomes new root

Search Algorithms

Range Query

FUNCTION List RQ(q, ϵ):

List C // candidates (MBRs/Objects)

List Result

C.insert(root)

WHILE(not C.isEmpty())

 E := C.removeFirstElement()

 IF E.isMBR()

 FOREACH $F \in E.children()$

 IF $\minDist(F, q) < \epsilon$

 C.insert(F)

 ELSE

 Result.insert(E)

RETURN Result

Nearest Neighbor Query (Top-Down Best-First-Search)

FUNCTION Object NNQuery(q):

PriorityQueue Q // sorted by mindist

Q.insert(0, root)

WHILE(not Q.isEmpty())

 E := Q.removeFirstElement()

 IF E.isMBR()

 FOREACH $F \in E.children()$

 Q.insert($\minDist(F, q)$, F)

 ELSE

 RETURN E

Spatial Joins

ϵ -Range Join Definition

$$S \bowtie_{\epsilon} G = \{(g, s) \in G \times S \mid \text{dist}(g, s) \leq \epsilon\}$$

Allows determining objects within AoI for each entity with single operation.

R-tree Spatial Join Algorithm


```

FUNCTION rTreeSimJoin(R, S, result,  $\epsilon$ )
  IF R.isDirectoryPage() or S.isDirectoryPage()
    FOREACH  $r \in R.children()$ 
      FOREACH  $s \in S.children()$ 
        IF  $\minDist(r,s) \leq \epsilon$ 
          rTreeSimJoin( $r,s,result,\epsilon$ )
  ELSE // R,S are data pages
    FOREACH  $p \in R.points$ 
      FOREACH  $q \in S.points$ 
        IF  $dist(p,q) \leq \epsilon$ 
          result.insertPair( $p,q$ )
  RETURN result

```

Problems of Data Volatility

Issues with Spatial Movement

- In games, majority of objects move several times per second
- **Changing position by delete/insert:**
 - Dynamic changes negatively influence structures
 - Causes overhead (queries, inserts, under/overflow handling)
- **Changing position via dedicated operations:**
 - Page region expansion increases overlap
 - Moving objects between regions affects balance

Conclusion: Dynamic calculation either has huge computational overhead or degenerates data structures.

Throw-Away Indices

Concept

For highly volatile data, rebuilding with bulk load is more efficient than updating existing structures.

Implementation

- Use 2 index structures:
 - **I₁:** Represents positions from last consistent tick (used for queries)
 - **I₂:** Created simultaneously
 - Via Bulk-Load: Less concurrency, fast creation, good structure
 - Dynamic creation: Higher effort, potentially worse structure
- At new tick start: I₂ used for queries, I₁ deleted and rebuilt

Key principle: Use tree only if creation + query processing time < brute force query time.

Design Considerations

Spatial problems depend heavily on game design:

- Number and distribution of spatial objects
 - Number and distribution of players
 - Environmental model (fields, 2D, 3D)
 - Note: 3D environment doesn't necessitate 3D indexing
 - Movement type and speed of objects
-

2. Virtual Environments

Overview

This section covers:

- Structure of Games
- Game States and Game Entities
- Transitions and Temporal Models
- Games and Simulations
- Games as Stochastic Processes

Building Blocks in Game Architecture

Architecture Layers (top to bottom)

1. Game Content

2. Game Engine

- Sound Manager
- Animation Engine
- Rendering Engine
- Graphics Engine
- Game Core
- Persistence Layer
- Script Language
- User Interface
- AI Engine

- Physics Engine
- Network Module

3. **Hardware Abstraction Layer** (DirectX, OpenGL, etc.)

4. **Hardware** (Sound Card, Graphics Card, Input devices, Network)

AI-Relevant Components

Directly involved:

- AI Engine: computes actions
- Game Core: represents the environment
- Physics Engine (sometimes)

Relevant infrastructure:

- Network layer: aspects of distributed environments
- Spatial management (in core or physics engine)
- Persistence layer: store data to analyze and learn from
- Scalability: required for massive multiplayer environments or fast forward simulations

Game States

Definition

All data representing the current state of the game:

- Objects, attributes, relationships (think ER or UML models)
- Contains all dynamic information
- Lists all game entities
- Contains all attributes of game entities
- Information concerning the whole game

Not necessarily in game state:

- Static information
- Environmental models/maps
- Preset attributes of game entities

Good Design Pattern

Model-View-Controller (MVC) Pattern:

- **Server:** Managing game state / no visualization (model and modification logic)

- **Client:** Only parts of game state but requires I/O and visualization (view and controller)

Example: Chess Game State

- **Game information:**
 - Players and assigned colors (black/white)
 - Game mode (with/without chess clock)
- **Game state:**
 - Positions of all figures / field occupation
 - Next player to move
 - Time left for both players (if using chess clock)

Transitions

What triggers game state changes?

- Player actions (e.g., chess move)
- Agent actions (AI move)
- Transition models (e.g., falling objects)

Need for synchronization:

- Model describing game time
- Organize temporal order of changes
- Restrict actions per time interval
- Ensure fair opportunities for all players
- Synchronize with wall-clock time

Consistency and Transitions

- Given consistent game state
- Transitions change valid state to new valid state
- Implement the rules of the game

Model-based Transitions Example

Ball on slope:

1. Environmental model decides ball rolls down
2. Transitions change position and state (acceleration)

Modeling Transitions

- Validity depends on game rules
- Distinguish between:
 - Computing an action (AI)
 - Environment's reaction (transition model)
- Transition models either:
 - Prevent prohibited actions
 - Compute valid environmental reaction

Examples:

- **Chess:** Only allow possible moves
- **Racing game:** Collisions, acceleration, deformation, disentanglement (physics engine)

Physics Engines

Implementation

- Solid state physics and classical mechanics
- Game entity provides parameters:
 - Spatial extension (polygon mesh, simplified: cylinders, MBRs)
 - Movement vectors
 - Mass
 - Other physical properties

Computational Requirements

- Uses differential equations
- Realistic effects require:
 - Large tick rates
 - Detailed models
- Large computational effort:
 - Precomputation
 - Numerical approximations

Detail Level vs. Performance Trade-off

- Transition models often use rough approximations
- More detail = less extended environment possible
 - Fewer entities
 - Lower framerates

- Detail level depends on game concept

Example contrast:

- **FPS:** Models physics of flying bullets
- **RPG:** Entire process modeled as random process

Temporal Models

Requirements

- Proper timing for transitions
- Time intervals for physics calculations
- Turn enforcement for board games
- Handle simultaneous actions
- Organize and synchronize transitions

Three Main Temporal Models

1. Turn-based Model

- **Process:**
 - Agents queried in fixed order
 - Actions sequential or parallel
 - Transitions computed in fixed order
 - Triggered when agents complete moves
- **Characteristics:**
 - Game state fixed during decision
 - Agents can block progression
 - Action order/amount determined by model
- **Pros:**
 - Clear, understandable temporal management
 - Time progresses based on slowest agent
- **Cons:**
 - Not suitable for real-time
 - No advantage for faster agents

Example: Round-based RPG

1. Environment randomly decides attack order

2. Agents choose actions sequentially
3. Environment computes state after each action
4. Environment progresses transition model

2. Transaction Model

- **Process:**
 - Actions queued and processed FIFO
 - Each transition processed sequentially
 - Game state = result of last transition
 - Agents submit actions at any speed
- **Essentially:** ACID database with sequential transactions
- **Pros:**
 - Guarantees valid game state
 - Faster agents can submit more actions
- **Cons:**
 - No sync with wall-clock time
 - No simultaneous actions

3. Tick System (Soft Real-Time)

- **Process:**
 - Transitions at fixed time intervals
 - Ticks represent game time intervals
 - All actions within tick treated as simultaneous
 - Joint state transition computation
- **Characteristics:**
 - Real-time w.r.t. tick length
 - Action delayed at most one tick
 - "Soft" = processing can exceed wall-clock interval
- **Pros:**
 - Synchronizes game and wall-clock time
 - Fair action rules per time unit
 - Concurrency possible
- **Cons:**
 - Handling lags

- Conflict resolution needed
- Chronological order issues

Actions vs. Transactions

Concurrent actions must be computed interdependently (not serializable)

Example scenario:

- Character A has 100/100 HP
- At time t:
 - A suffers 100 HP damage from B
 - A receives 100 HP healing from C

Outcomes:

- **Isolation (healing first):** A dies (overheal → damage)
- **Isolation (damage first):** A dies (can't heal dead)
- **Concurrent:** Effects cancel, A survives

Conflicts During Concurrency

Example: Two players pick up gold coin simultaneously

Conflict Resolution Options:

1. **Delete both actions:** Conflict detection + possible reset
2. **Random pick:** One succeeds, other fails
3. **First action executes:** Natural order (may not be fair)

Consistency During Action Handling

Problem: Reading already changed data

Solutions:

1. **Shadow Memory:**
 - Two game states: G_1 (active/consistent) and G_2 (inactive/changing)
 - All reads from G_1 , all writes to G_2
 - On tick completion: swap active/inactive
2. **Fixed sequence of read/write operations:**
 - Break down action components
 - Rearrange operations

- Handle all actions simultaneously

Time Model Extensions

Turn-based can extend to tick-system by:

- Collecting actions from all agents each turn
- Setting time-out for acting
- Computing transitions regardless of submissions
- Providing game time per tick to environment model

Comparison to Scientific Simulations

Similarities

- Usually use tick systems
- Often no actors, just physics engine
- Agent-based simulations use AI
- Need tick systems for fixed real-world time interpretation

Differences

- Simulations don't require wall-clock sync
- Technically similar but transition models target maximal realism

Event-based vs. Tick-based Simulations

Event-based applicable if:

- Can schedule when each agent acts next
- Can compute future environment state without agent actions

Process:

1. Determine next tick for each agent
 2. Organize in priority queue
 3. Progress environment to top tick
 4. Collect all actions for that tick
- **Benefit:** Skip eventless ticks for efficiency

Games as Stochastic Processes

Rationale

- Mathematical model helpful for analysis
- Games involve non-deterministic elements:
 - Environmental random components
 - Player action variation
- Otherwise all matches identical
- Match = stochastic process

Discrete Time Homogeneous Markov Process

Definition requires:

- Set of states S
- Stochastic transition function $T: P(s'|s)$ for all $s, s' \in S$

Optional:

- Start distribution $B: P(s'| -)$ for first state
- Terminal states S_{terminal} ending process

Example calculation:

$$\begin{aligned} P(ACBB) &= P(A|-) \cdot P(C|A) \cdot P(B|C) \cdot P(B|B) \cdot P(-|B) \\ &= 0.3 \cdot 0.15 \cdot 0.7 \cdot 0.4 \cdot 0.1 \end{aligned}$$

Properties of Markov Processes

1. First-order Markov Property:

- Process is memoryless
- $s_{\{t+1\}}$ depends only on s_t

2. Discrete time:

- Process runs in discrete steps
- Corresponds to turn/tick-based time

3. Continuous time Markov chains:

- Real-time transactions
- Need to predict next transition time

Homogeneous vs. Inhomogeneous Processes

Homogeneous:

- Same transition function each time step
- Example: PageRank algorithm

Inhomogeneous:

- Transition function varies over time
- $p(s_{t+1}=A|s_t=B) \neq p(s_{t+i+1}=A|s_{t+i}=B)$ for $i \geq 1$
- Example: More mobs spawn at night

Markov Processes and Games

Games usually fulfill Markov Property:

- Game state contains all necessary information
- Player actions subsumed into transition function
- Without randomness/unknown actions → deterministic

Characteristics:

- Discrete time = tick/turn-based time
- Often inhomogeneous (player behavior changes)
- Extract agent actions later (Markov Decision Processes)
- Environment often homogeneous over time

Summary

Key Concepts

1. Spatial Management:

- Critical for game scalability
- Various techniques: zoning, indexing, publish-subscribe
- Trade-off: update/query processing time

2. Game Architecture:

- Separation of concerns (MVC pattern)
- AI integrates with core systems
- Physics and spatial management crucial

3. Temporal Models:

- Turn-based, transaction, tick systems
- Each has specific use cases
- Synchronization challenges

4. Stochastic Modeling:

- Games as Markov processes
- Basis for data modeling

- Transition functions describe mechanics and behavior

5. **Performance Considerations:**

- Spatial queries grow quadratically
- Index structures must justify overhead
- Throw-away indices for volatile data

These concepts form the foundation for understanding and implementing AI systems in modern games, balancing performance, scalability, and gameplay requirements.